

Forget ODBC! Here's a New Db2 Connector

Jonathan Zak, IBM

Software Developer

jonathan.zak@ibm.com

Sanjula Ganepola, IBM

Software Developer

sanjula.ganepola@ibm.com

Agenda

- Maepire Overview
- Architecture and Core Tenets
- Comparisons versus JDBC and ODBC
- Deep dive into Node.js client SDK
- Demo

What is Mapepire?

Welcome to Mapepire

A cloud-friendly IBM i database access layer, built with simplicity and performance in-mind.

Find out more →

Pick your client language ⓘ



*Super easy to use way to access
Db2 for i from any application*



Mapepire Origin Story...

January 2020

- VSCode "Code for IBM i" extension includes basic Db2 support



February 2022

- Work begins on Server component to power Db2 features in VSCode



March 2022

- First release of VSCode Db2 for i extension

July 2023

- VSCode Db2 for i extension publishes v0.3.0, the first release leveraging server component (v0.3.0)

August 2024

- Mapepire is born!



Server Component

SDK architecture

Python

Java

TypeScript

C#

PHP


FUTURE

How does it work?

Connect to Database

```
{  
  "id": "id2",  
  "type": "connect",  
  "technique": "tcp",  
  "application": "Java client"  
}
```

Communication over
secure web sockets



Application uses
Mapepire client SDK

```
{  
  "id": "id2",  
  "job": "112480/QUSER/QZDASOINIT",  
  "success": true  
}
```

Managed using
Service Commander

How does it work?

Query the Database

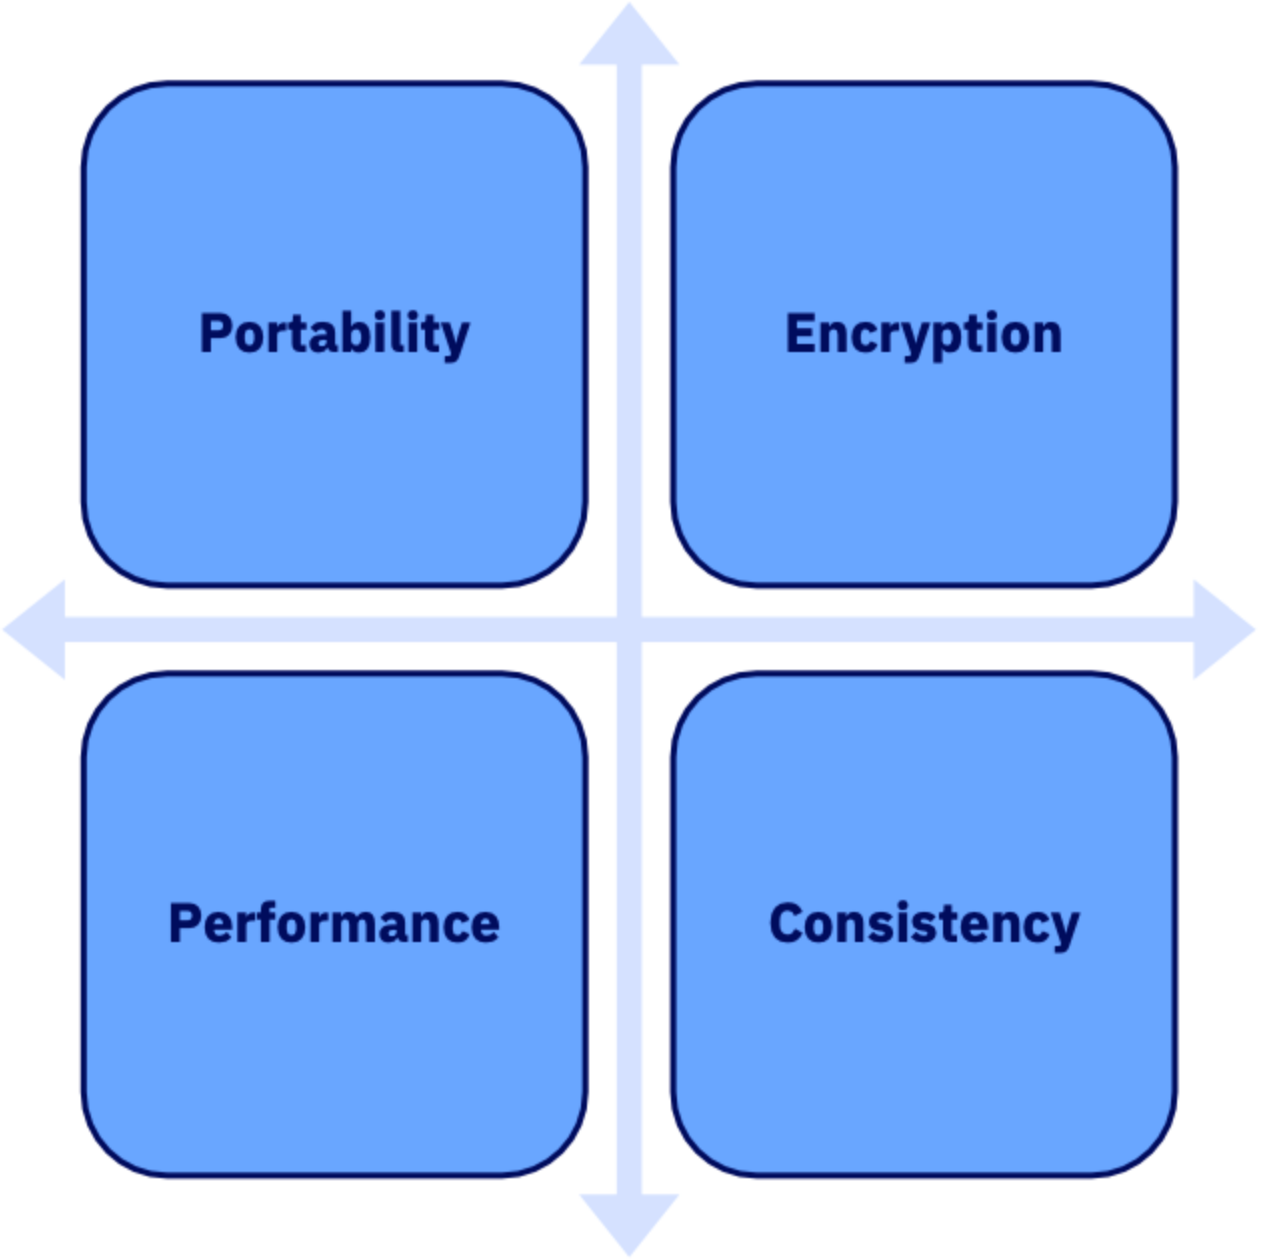
Request

```
{
  "id": "query3",
  "type": "sql",
  "sql": "SELECT DEPTNO FROM SAMPLE.DEPARTMENT
WHERE DEPTNO = 'A00'",
  "terse": false,
  "rows": 100
}
```

Response

```
{
  "id": "query3",
  "has_results": true,
  "update_count": -1,
  "metadata": {
    "column_count": 1,
    "job": "112480/QUSER/QZDASOINIT",
    "columns": [
      {
        "name": "DEPTNO",
        "type": "CHAR",
        "display_size": 3,
        "label": "DEPTNO",
        "precision": 3,
        "scale": 0
      }
    ]
  },
  "data": [
    {
      "DEPTNO": "A00"
    }
  ],
  "is_done": true,
  "success": true
}
```

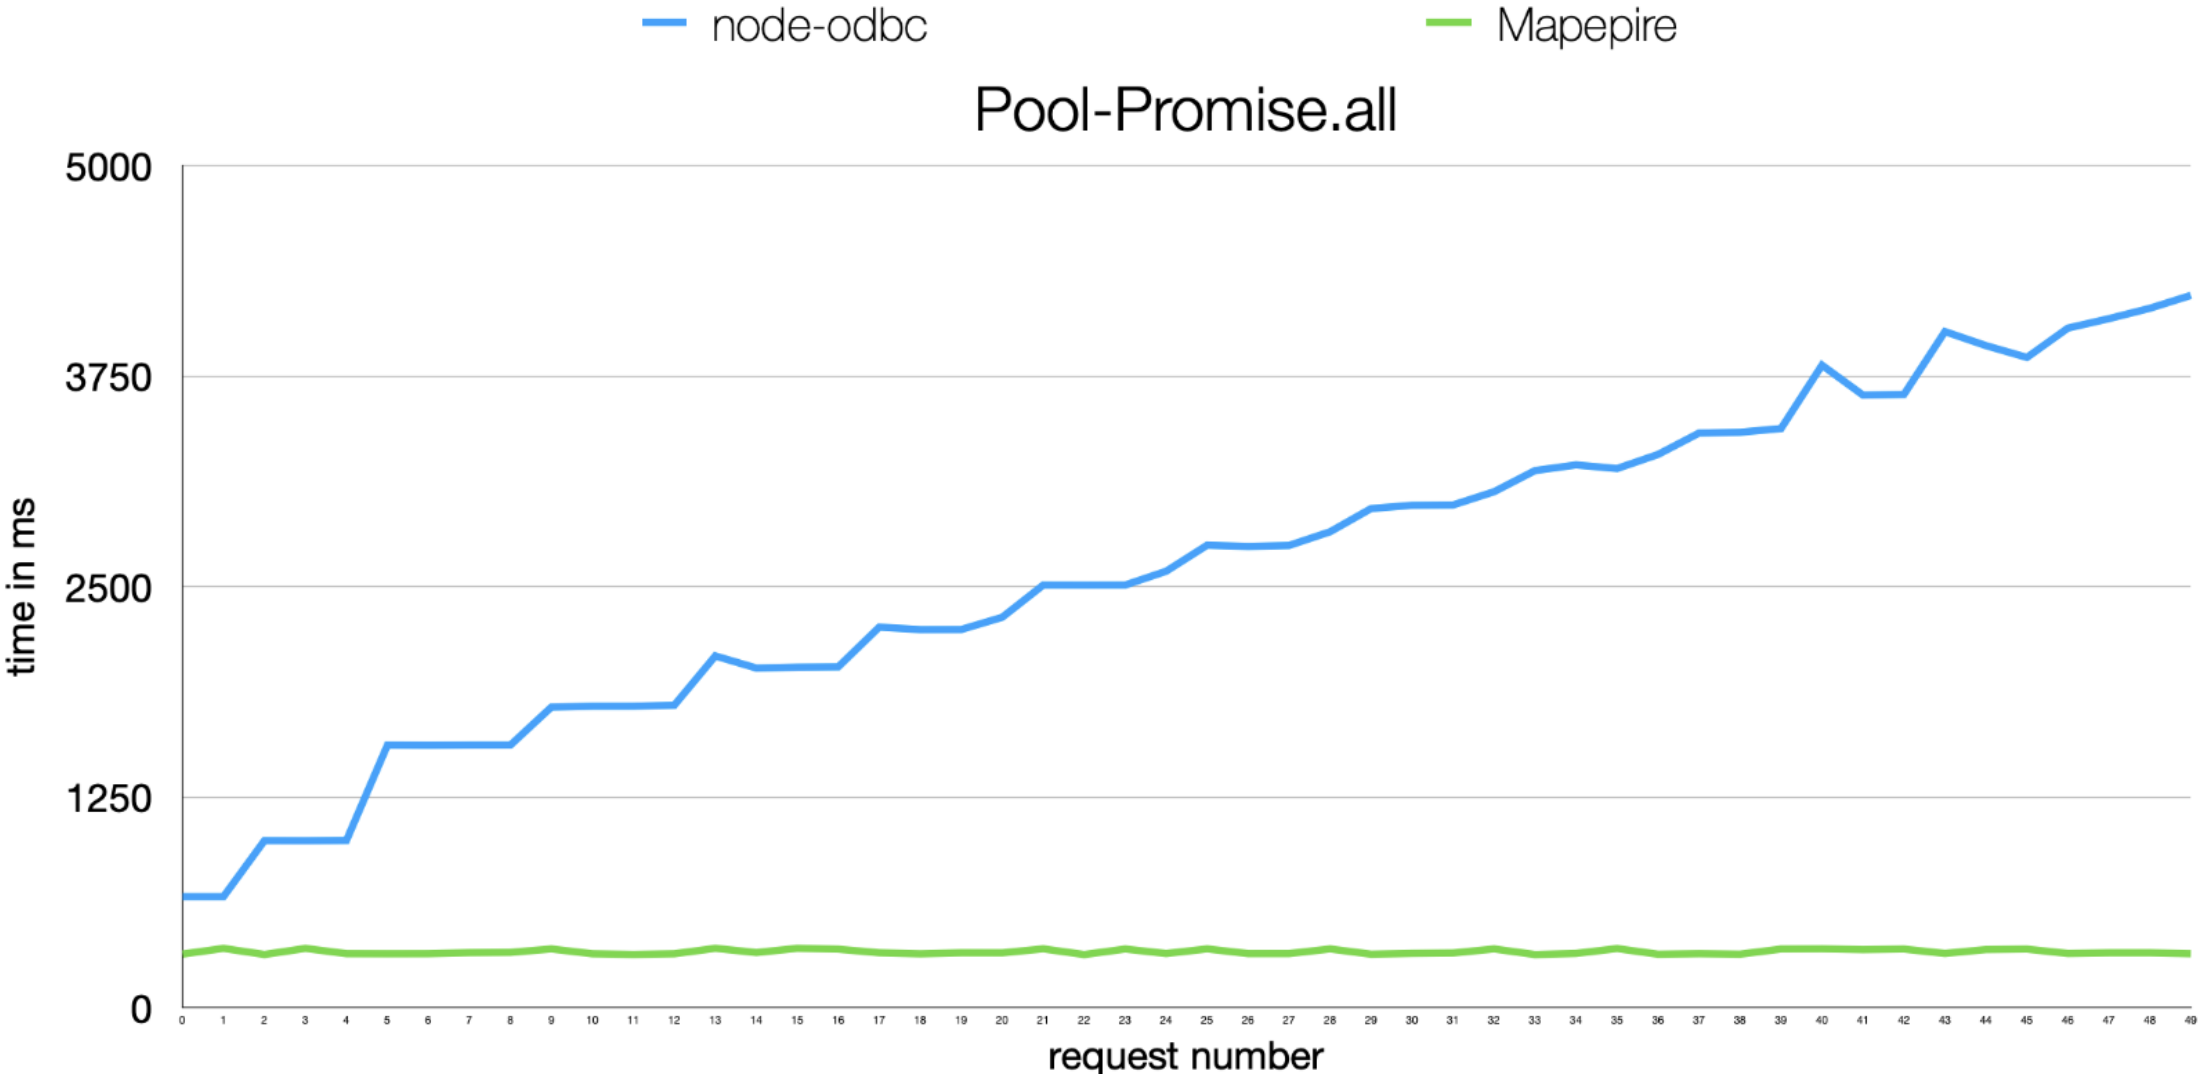
Core Tenets



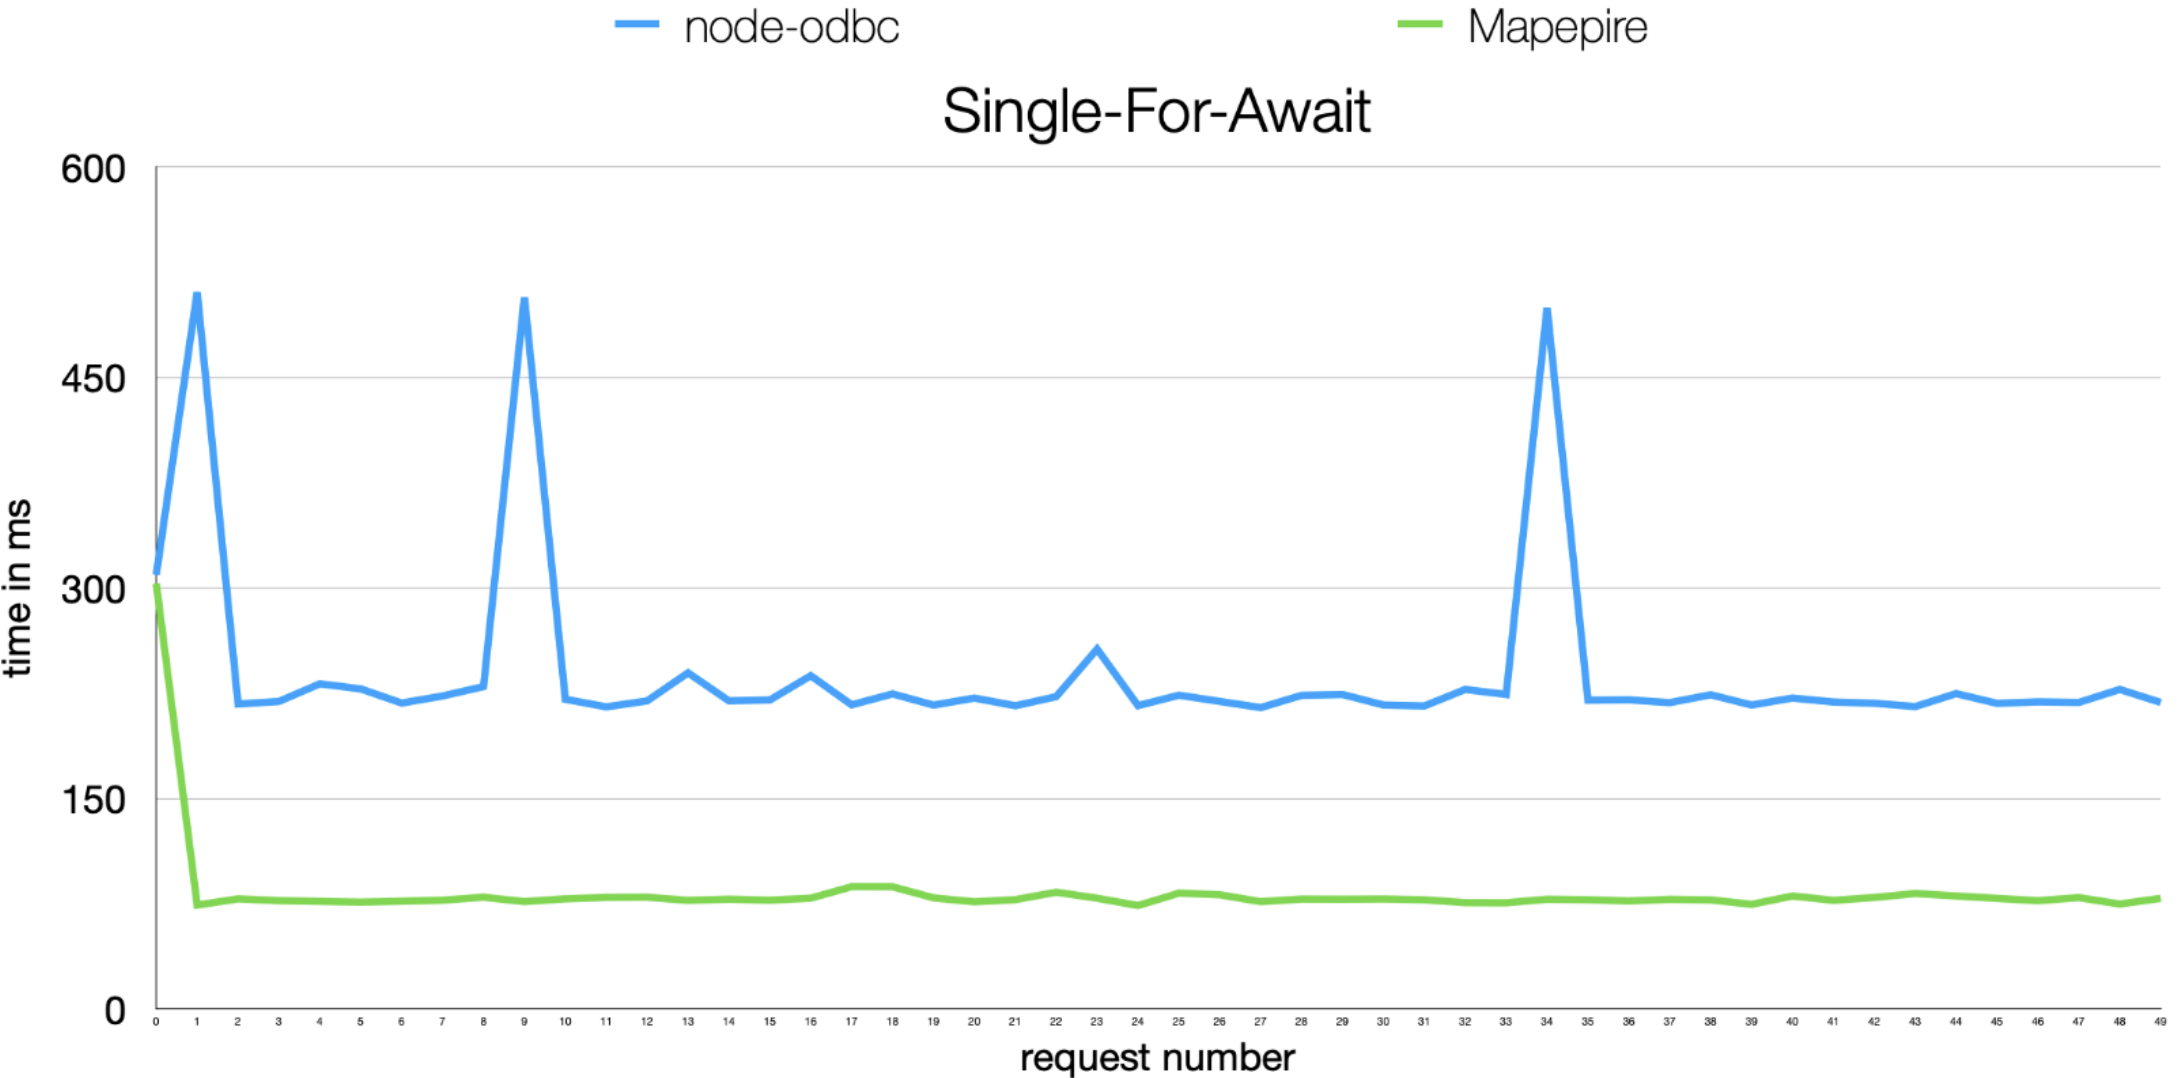
The biggest benefit of Mapepire.... Portability!!

	JDBC	ODBC	Mapepire
Runs in WatsonX.ai Jupyter notebooks	✗	✗	✓
Runs in Rocket AI Hub programmer portal	✗	✗	✓
Runs in Rocket Cognitive Environment	✓*	✗	✓
Runs in Alpine Linux containers	✓	✗	✓
Runs in Raspberry Pi	✓	✗	✓
Runs in Arduino	✗	✗	✓

Some performance comparisons



Some performance comparisons



How to encrypt data with JDBC/ODBC

1. Log into DCM
2. Create a local certificate authority (CA) store
3. Create a local CA certificate
4. Record the value of the auto created CA label
5. Create the *SYSTEM certificate store (if needed)
6. Create a new server certificate
7. Sign the server certificate with your local CA
8. Assign new server certificate to host server applications
9. Restart Host Servers
10. On client, download the server's certificate authority to a local trust store (or configure TLS to ignore completely)

How to encrypt data with Mapepire?

Step 1

- Use Mapepire

Step 2

- Go to bed

How does Maepire make it so easy?

Option 1: Custom certificate

- Admin explicitly defined a custom certificate by configuring a certificate store:
 - File name: `/QOpenSys/etc/maepire/cert/server.jks`
 - Format: `JKS`
 - Store Password: `maepire`
 - Key Password: `maepire`
 - Certificate Alias: `maepire`
- Check out documentation for full instructions: <https://maepire-ibmi.github.io/guides/sysadmin/>

Option 2: Let's Encrypt

- Use Let's Encrypt (ex. generated by CertBot)
- Maepire server will automatically use it as the server certificate
- Certificate must exist in the following location used by CertBot:
`/etc/letsencrypt/live/<hostname>`

Option 3: Self-signed certificate

- If no certificate, the server automatically generates its own self-signed certificate

What does TLS provide?

Encryption

- Data isn't sent "in the clear"

Authentication

- Client ensures the server certificate is valid
- Client ensures the server certificate is signed by a trusted authority
- Client checks that the hostname matches that of the certificate

Mapepire's back-end is JDBC

- Mapepire is an interface in front of JTOpen and JDBC
- Mapepire utilizes QZDASOINIT or, more likely, QZDASSINIT (S = "secure") jobs
- All considerations for ODBC/JDBC server job scalability and security still apply
 - Object authority still applies
 - Any ODBC/JDBC exit points will still work to control traffic and access

How to manage the JDBC workload?

- By default, all QZDASOINIT/ QZDASSINIT jobs run in QUSRWRK
- Questions:
 - How to control out-of-control queries from query tools?
 - How to know which application is using up resources?
 - How to let critical users get the resources they need while not letting long queries take over the system?
 - How to manage your JDBC jobs more effectively?

Separate jobs by application, user, etc.

- Configure your QZDASSINIT jobs to run in separate subsystems, based on your criteria
 - JDBC SHOP
 - JDBC ADHOC
 - JDBC NODE
- Then performance "waits" can be aggregated by subsystem and you can configure memory, etc. per subsystem
- Easier troubleshooting as JDBC jobs from different applications will not interact
- Also limit with these techniques <https://www.ibm.com/support/pages/setting-limitations-resources-used-qzdasoinit-prestart-jobs>
- Details: <https://www.seidengroup.com/2022/05/04/simplify-with-subsystems/>

Configure JDBC Prestart jobs



- Configure the right number of prestart jobs
 - ODBC/JDBC prestart jobs are QZDASO(S)INIT in QUSRWRK
- Check out your current configuration:
 - [DSPSBSD SBSD\(QUSRWRK\)](#)
 - Choose 10, Prestart job entries
 - Type 5 next to QZDASOINIT

```
Display Prestart Job Entry Detail
Subsystem description:  QUSRWRK      Status:  ACTIVE      System:  SV12
Program . . . . . :  QZDASOINIT
Library . . . . . :  QSYS
User profile . . . . . :  QUSER
Job . . . . . :  QZDASOINIT
Job description . . . . . :  QDFTSVR
Library . . . . . :  QGPL
Start jobs . . . . . :  *YES
Initial number of jobs . . . . . :  1
Threshold . . . . . :  1
Additional number of jobs . . . . . :  2
Maximum number of jobs . . . . . :  *NOMAX
Maximum number of uses . . . . . :  200
Wait for job . . . . . :  *YES
Pool identifier . . . . . :  1
```

Default ODBC/JDBC prestart job settings



Low defaults

- Initial jobs = 1, threshold = 1, additional jobs = 2
- Change as needed:
CHGPJE SBSDB(QSYS/QUSRWRK) PGM(QSYS/QZDASOINIT)
STRJOBS(*YES) INLJOBS(xx) THRESHOLD(xx)
ADLJOBS(xx)
- How to determine **optimal values**? **DSPACTPJ** (coming up)

```
Display Prestart Job Entry Detail
System: SV12
Subsystem description: QUSRWRK      Status: ACTIVE
Program . . . . . : QZDASOINIT
Library . . . . . : QSYS
User profile . . . . . : QUSER
Job . . . . . : QZDASOINIT
Job description . . . . . : QDFTSVR
Library . . . . . : QGPL
Start jobs . . . . . : *YES
Initial number of jobs . . . . . : 1
Threshold . . . . . : 1
Additional number of jobs . . . . . : 2
Maximum number of jobs . . . . . : *NOMAX
Maximum number of uses . . . . . : 200
Wait for job . . . . . : *YES
Pool identifier . . . . . : 1
```

How many jobs are needed?

- DSPACTPJ SBS(QUSRWRK) PGM(QZDASOINIT)
- More details: <https://www.ibm.com/docs/en/i/7.4?topic=jobs-tuning-prestart-job-entries>

```

Prestart jobs:
-
  Current number . . . . . : 29
  Average number . . . . . : 21.9
  Peak number . . . . . : 49

Prestart jobs in use:
  Current number . . . . . : 27
  Average number . . . . . : 18.4
  Peak number . . . . . : 46
```

```

Program start requests:
-
  Current number waiting . . . . . : 0
  Average number waiting . . . . . : .0
  Peak number waiting . . . . . : 6
  Average wait time . . . . . : 00:00:00.0
  Number accepted . . . . . : 100299
  Number rejected . . . . . : 0
```

JDBC vs ODBC vs Mapepire

	JDBC	ODBC	Mapepire
Needs only a single port			✓
Data is always encrypted			✓
Manageable via system exit points	✓	✓	✓
Enhanced CCSID support	✓		✓
Runs in WatsonX.ai Jupyter notebooks			✓
Runs in lightweight containers (for instance Alpine Linux)	✓		✓
Directly supports multiple client languages			✓

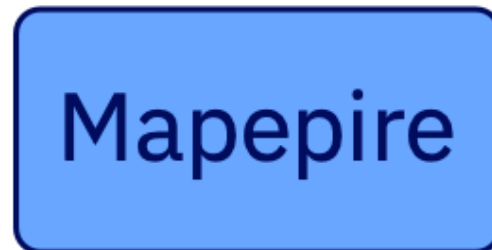
Single port? Big deal!

TCP distance to first database operation

JDBC/ODBC



Maepire



Distinct TCP flows for a JDBC program!!

```
1  try (AS400 hi = new AS400("myhostname", "uid".toCharArray(), "password".toCharArray())) {
2      AS400JDBCDataSource ds = new AS400JDBCDataSource(hi);
3      Connection conn = ds.getConnection();
4      Statement s = conn.createStatement();
5      s.executeQuery("select * from QIWS.QCUSTCDT");
6
7      ResultSet rs = s.getResultSet();
8      while (rs.next()) {
9          System.out.println(rs.getString(1));
10     }
11     System.out.println("done");
12 }
```

Meanwhile....

✨ **Maepire only needs 2 TCP flows**

1. Connect and allocate a job
2. Run a query

Host Server	
1::S - 7003	Exchange Client/Server Attributes
1::R - F003	Exchange Client/Server Attributes Reply
1::S - 7004	Retrieve Signon Information
1::R - F004	Retrieve Signon Information Reply
1::S - 7006	End Job Request
2::S - 7001	Exchange Random Seeds
2::R - F001	Exchange Random Seeds Reply
2::S - 7002	Start Server
2::R - F002	Start Server Reply
2::S - 1F80	Set Attributes
2::R - 2800	SQL Requested Data Returned
2::S - 1D00	Create and init RPB with no based-on RPB
2::S - 1803	Prepare/Describe
2::R - 2800	SQL Requested Data Returned
2::S - 180E	Open/Describe/Fetch
2::R - 2800	SQL Requested Data Returned

How to get started with Mapepire?

- Install the mapepire client with:
`npm install @ibm/mapepire-js`
- Create a .env file specifying:
 - HOST
 - USER
 - PASSWORD
 - PORT (this is the Mapepire server port)

```
⚙️ .env
1  HOST="myhost.somewhere.com"
2  PORT="8076"
3  USER="jimbob"
4  PASSWORD="letmein4"
5
```

```
const creds = {
  host: process.env.HOST,
  port: process.env.PORT,
  user: process.env.USER,
  password: process.env.PASSWORD,
};

const ca = await mapepire.getCertificate(creds);
creds.ca = ca.raw;
const job = new mapepire.SQLJob();
await job.connect(creds);
```

Launch the Mapepire server

- Install the Mapepire server component: `yum install mapepire-server`
- Install Service Commander: `yum install service-commander`
- Launch mapepire: `sc start mapepire`

```
cd /opt/mapepire/lib/mapepire
```

```
## Start
```

```
sc start mapepire.yaml
```

```
## Check it's running
```


```
sc check mapepire.yaml
```

```
## Stop
```

```
sc stop mapepire.yaml
```

```
## Check it's stopped
```

```
sc check mapepire.yaml
```

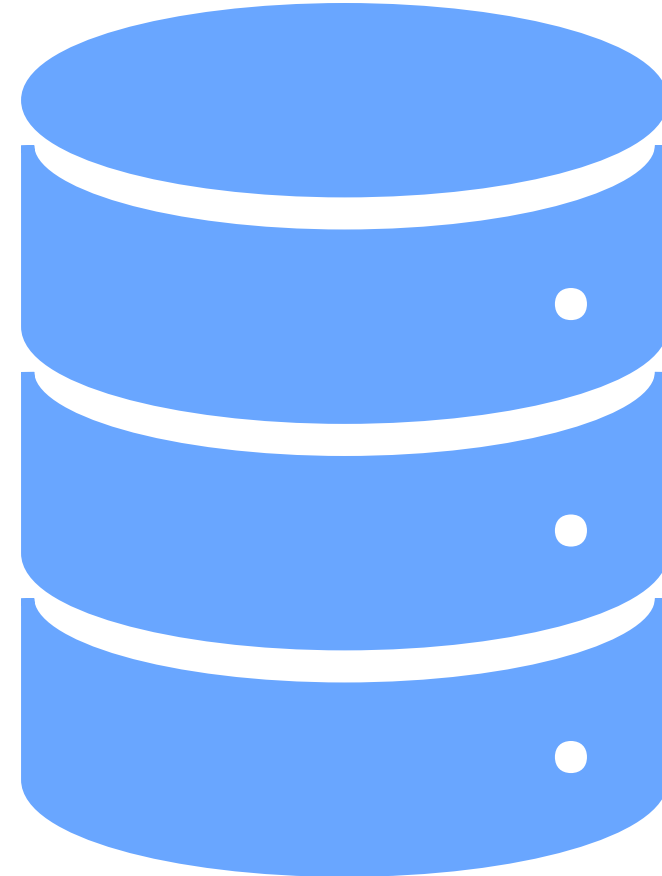


```
-bash-5.2$ sc start mapepire  
Performing operation 'START' on service 'mapepire'  
Service 'Mapepire Server' successfully started
```

Now it's time to write some queries!

Here is some of the functionality:

- Connect to the database
- Run CL commands
- Run SQL commands
- Run SQL commands as a batch
- Run procedures
- Explain a statement
- Get server job
- Get Mapepire version
- Check liveness
- Set config
- Get trace data
- Close a connection



Connecting to the database

Use the connect method of the [SQLJob](#) class.

```
// Connect to the database  
const creds = ENV_CREDS ;  
const job = new mapepire.SQLJob() ;  
await job.connect(creds) ;
```

Get the server certificate

Use the `getCertificate` method to retrieve the TLS certificate from the db2 server.

If the server is using a self-signed certificate, you can add it to your connection credentials to specify you trust this certificate.

```
const creds = ENV_CREDS ;  
const ca = await mapepire.getCertificate(creds);  
creds.ca = ca.raw;
```

Executing a query

Use the `query` method of `SQLJob` to construct a query.

Then call the queries `execute` method.

```
await job.connect(creds);
const query = await job.query<any>("select * from sample.department");
const res = await query.execute();
await query.close();
await job.close();
```

```
✓ data = (13) [{...}, {...}, {...}, {...}, {...}, {...}, {...}, {...}, {...}, {...},
> 0 = {DEPTNO: 'A00', DEPTNAME: 'SPIFFY COMPUTER SERVICE DIV.',
> 1 = {DEPTNO: 'B01', DEPTNAME: 'PLANNING', MGRNO: '000020', AD
> 2 = {DEPTNO: 'C01', DEPTNAME: 'INFORMATION CENTER', MGRNO: '0
> 3 = {DEPTNO: 'D01', DEPTNAME: 'DEVELOPMENT CENTER', MGRNO: nu
> 4 = {DEPTNO: 'D11', DEPTNAME: 'MANUFACTURING SYSTEMS', MGRNO:
> 5 = {DEPTNO: 'D21', DEPTNAME: 'ADMINISTRATION SYSTEMS', MGRNO
> 6 = {DEPTNO: 'E01', DEPTNAME: 'SUPPORT SERVICES', MGRNO: '000
```

Executing a procedure

You can execute a procedure the same way that you would execute an SQL statement!

```
const job = new SQLJob();
await job.connect(creds);

const testProc = `
create or replace procedure ${TEST_SCHEMA}.procedure_test_char(
  in p1 char(5),
  inout p2 char(6),
  out p3 char(7)
)
BEGIN
  set p3 = rtrim(p1) concat rtrim(p2);
  set p2 = '';
END
`;

const queryA = job.query<any[]>(testProc);
await queryA.execute();
await queryA.close();
```

Executing a CL command

Use the CL method of the [SQLJob](#) class to execute CL commands.

```
const job = new SQLJob();  
await job.connect(creds);  
const query = await job.clcommand("WRKACTJOB");  
const res = await query.execute();  
await job.close();
```


Executing a prepared statement

Use the `query` method of `SQLJob` to construct a query

Add the parameters for the prepared query

Then call the queries `execute` method

```
const job = new SQLJob();
await job.connect(creds);
const query = await job.query<any>(
  "SELECT * FROM SAMPLE.SYSCOLUMNS WHERE COLUMN_NAME = ?",
  {
    parameters: ["Value"],
  }
);
const res = await query.execute();
await query.close();
await job.close();
```

Explain an SQL statement

Use the [explain](#) method of [SQLJob](#) to explain a SQL statement and return the result.

```
const job = new mapepire.SQLJob();  
await job.explain("select * from SAMPLE");  
await job.close();
```

Executing a batch of statements

Use the `query` method of `SQLJob` to construct a prepared statement

Use a 2D array of parameters to ensure the query is constructed as a batch operation

Then call the queries `execute` method

```
let query = job.query("INSERT INTO SAMPLE.EMPLOYEES values (?, ?)", {  
  parameters: [  
    ["SANJULA", "416 345 0879"],  
    ["TONGKUN", "647 345 0879"],  
    ["KATHERINE", "905 345 1879"],  
    ["IRFAN", "647 345 0879"]  
  ],  
});  
let res = await query.execute();
```

Executing statements using a pool

Initialize a new `pool` object

Call the pool's `execute` method to run a query using any free job

Call the pool's `end` method to destroy the pool and cleanup resources

```
let pool = new Pool({ creds, maxSize: 5, startingSize: 5 });
await pool.init();
// Initiate a bunch of jobs
const executedPromises = [
  pool.execute("select * FROM SAMPLE.SYSCOLUMNS"),
  pool.execute("select * FROM SAMPLE.SYSCOLUMNS"),
  pool.execute("select * FROM SAMPLE.SYSCOLUMNS"),
];
const res = await Promise.all(executedPromises);
await pool.end();
```

Retrieve the SQL Job that is running the query

Use the `getHostJob` method of `Query` class to return the `SQLJob` that is running the query

```
const job = new SQLJob();  
await job.connect(creds);  
const query = await job.query<any>("select * from sample.department");  
const res = await query.execute();  
  
const queryJob = query.getHostJob();
```

Retrieve the status for the SQL job

Use the `getStatus` method of `SQLJob` to retrieve the status of the job

Can be any of `notStarted`, `connecting`, `ready`, `busy`, `ended`

```
const job = new mapepire.SQLJob();  
await job.getStatus();  
...
```

Retrieve the state of the query

Use the `getState` method of `Query` class to return the state of the query

Can be one of "NOT_YET_RUN", "RUN_MORE_DATA_AVAILABLE", "RUN_DONE", "ERROR"

```
const job = new SQLJob();  
await job.connect(creds);  
const query = await job.query<any>("select * from sample.department");  
const res = await query.execute();  
const state = query.getState();
```

Retrieve the count of pending transactions

Use the `getPendingTransactions` method of `SQLJob` to retrieve the count of pending transactions.

```
const job = new mapepire.SQLJob();  
await job.getPendingTransactions()
```


Retrieve the number of ongoing requests for a job

Use the `getRunningCount` method of `SQLJob` to retrieve the number of ongoing requests for the job

```
const job = new mapepire.SQLJob();  
await job.getRunningCount()  
...
```

Retrieve the trace file path

Use the `getTraceFilePath` method of `SQLJob` to retrieve the file path of the trace file

```
const job = new mapepire.SQLJob();  
await job.getTraceFilePath()  
...
```

Set the backend log level

Use the `setTraceConfig` method of the `SQLJob` class to set the trace level and destination for the Mapepire server

In this case we are sending the trace data to a file, and we are using the most verbose level of logging

```
const job = new mapepire.SQLJob();  
await job.setTraceConfig("FILE", "DATASTREAM")
```

Retrieve the server trace data

Use the `getTraceData` method of `SQLJob` to retrieve the trace data from the backend

```
const job = new mapepire.SQLJob();  
await job.getTraceData()
```

Retrieve the Mapepire server version

Use the `getVersion` method of `SQLJob` to retrieve the Mapepire server version.

```
const job = new mapepire.SQLJob();  
await job.getVersion()
```

Consistent SDK behavior access languages

- Guided by a unified reference architecture
 - <https://mapepire-ibmi.github.io/reference/maintenance/referencearchitecture/>
- Similar experiences
 - Class names
 - Method names
 - Throwable types
 - Input parameters
 - Configuration options

Node.js vs Java Implementation

Node.js

```
// Initialize credentials
const creds: DaemonServer = { host: "HOST", port: 8076, user:
"USER", password: "PASSWORD", rejectUnauthorized: true, ca:
"CA" }
```

```
// Establish connection
const job = new SQLJob();
await job.connect(creds);
```

```
// Initialize and execute query
const query = job.query("SELECT * FROM SAMPLE.DEPARTMENT");
const result = await query.execute(3);
```

```
// Convert to JSON string and output
console.log(JSON.stringify(result));
```

Java

```
// Initialize credentials
DaemonServer creds = new DaemonServer("HOST", 8085, "USER",
"PASSWORD", true, "CA");
```

```
// Establish connection
SqlJob job = new SqlJob();
job.connect(creds).get();
```

```
// Initialize and execute query
Query query = job.query("SELECT * FROM SAMPLE.DEPARTMENT");
QueryResult<Object> result = query.execute(3).get();
```

```
// Convert to JSON string and output
ObjectMapper mapper = new ObjectMapper();
mapper.enable(SerializationFeature.INDENT_OUTPUT);
String jsonString = mapper.writeValueAsString(result);
System.out.println(jsonString);
```

Config

Great performance

Always encrypted

Flexibility

Any Hardware

Any Language



Demo

Any Questions?

Important Links

Mapepire

Documentation

<https://mapepire-ibmi.github.io/>

Server Component

<https://github.com/Mapepire-IBMi/mapepire-server>

Node.js Client

GitHub Repository

<https://github.com/Mapepire-IBMi/mapepire-js>

NPM

<https://www.npmjs.com/package/@ibm/mapepire-js>

Java Client

GitHub Repository

<https://github.com/Mapepire-IBMi/mapepire-java>

Maven Central

<https://central.sonatype.com/artifact/io.github.mapepire-ibmi/mapepire-sdk>

Python Client

GitHub Repository

<https://github.com/Mapepire-IBMi/mapepire-python>

PyPi


<https://pypi.org/project/mapepire-python/>

Service Commander

GitHub Repository

<https://github.com/ThePrez/ServiceCommander-IBMi>

For More Information

Links You Need	Twitter	#Hashtags
<p>IBM i Home Page: https://www.ibm.com/it-infrastructure/power/os/ibm-i (find link to Forrester Study and updated IBM i Strategy Whitepaper)</p> <p>IBM Strategy Whitepaper: https://www.ibm.com/it-infrastructure/us-en/resources/power/i-strategy-roadmap/</p> <p>IBM Client Success: https://www.ibm.com/it-infrastructure/us-en/resources/power/ibm-i-customer-stories/</p> <p>Support Life Cycle: https://www.ibm.com/support/lifecycle/</p> <p>License Topics: https://www-01.ibm.com/support/docview.wss?uid=nas8N1022087</p> <p>Fortra IBM i Marketplace Survey https://www.fortra.com/resources/guides/ibm-i-marketplace-survey-results</p>	<div></div> <div>@IBMSystems @COMMONug @IBMChampions @IBMSystemsISVs @IBMiMag @ITJungleNews @SAPonIBMi @SiDforIBMi</div>	<div>#PowerSystems #IBMi #IBMAIX #POWER9 #LinuxonPower #OpenPOWER #HANAonPower #ITinfrastructure #OpenSource #HybridCloud #BigData</div>

Forget ODBC! Here's a New Db2 Connector - Jonathan Zak and Sanjula Ganepola

Please take the last minute of this session to complete the evaluation. A direct link to the evaluation can be found using the QR code below.



